

Abstract

Machine Learning and Game Playing are two very hot topics. In this project we apply machine learning techniques to the board game Blokus. We picked Backpropagation Algorithm as our learning algorithm and applied it to a Neural Network with 7 layers. The project is implemented in python. We trained the Neural Network with over 1000 games. Significant improvement could be seen.

Introduction

Nowadays, Machine Learning becomes more and more popular. After the success of AlphaGo, the field catches the attentions all over the world. In 1959, Arthur Samuel defined machine learning as a “Field of study that gives computers the ability to learn without being explicitly programmed.”[] Developing traditional AI application without using machine learning strategy is the process that transfers human knowledge to formal computer algorithm and implements it in programming languages. This approach gives the computer the knowledge needed to complete the tasks rather than the abilities to learn the knowledge. Traditional AI approach without machine learning solves the Game Play problem typically in the following steps.

1. Given the current game state S_0 , calculate possible next game states $S_{next} = \{s_i | s_0 \xrightarrow{move_i} s_i\}$
2. Evaluate all game states in S_{next} with the Evaluation Function f . Find the best next game state S_{best} where $f(s_{best}) = \max\{f(s_i) | s_i \in S_{next}\}$ and corresponding move m_{best} . ($s_0 \xrightarrow{move_{best}} s_{best}$)
3. Apply $move_{best}$.

Obviously the key element in Game Play AI application is the Evaluation Function. Designing Evaluation Function is the process to transfer human knowledge or heuristics into a formula that gives a meaningful Evaluation Score to help with decision making.

The problem with the traditional AI

1. The Evaluation Function is hard to design. There are too many factors needed to be considered in order to give a meaningful score.
2. Typically the Evaluation Function only considers the current game state without looking at the future, which makes the Evaluation Score no so meaningful. Min-Max search[] was introduced to solve the problem. But as the number of game states grows exponentially, this approach becomes not practical.

Machine Learning overcomes the problem by giving computer the abilities to learning how to evaluate the game states. We use Artificial Neural Network as the model of Machine Learning. A game state is converted to a input vector which gets plugged into the Neural Network. The output of the Network is

the Evaluation Score of the game state. We use back-propagation algorithm as our learning algorithm. In the next section, we will talk about our approach in details.

Backpropagation Details

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Implementation Details

When we were assigned the task of writing a blokus-learning algorithm, our first response was to break our task into pieces:

- Some sort of host-program, game-engine, something that could manage the game board, keep track of it's current state, draw it on the screen, keep track of the players' pieces, and calculate possible moves for the players.
- A function that would 'act' for the player, it would somehow take in the current board, and decide what piece to place. This function would be called by the game-engine each turn.
- Some kind of "learning"-method, it would hopefully change the way the 'act'-ing-function would operate.

Before we considered the exact implementation of these pieces, we knew that we could make our lives much easier if we didn't start of scratch, and instead build off of other developers's code. That's when we discovered a repository on github that already had developed the game-engine, players, and simple user-interface. It out-of-the box ran the user-player against a random-computer-player, and drew the board on an ASCII grid printed out on the terminal. We thought this was a perfect starting point, so we began exploring through the code, to find develop a basic idea of the flow of the program, and to start getting an idea of how we can change this program to suit our needs. Each player was an object containing an array of pieces, a score, a name, corner to start from, and a function pointer. Each player was given a function-pointer to call when it needed to decide what piece to place where (the 'act'-ing function). The player objects are handed their 'act'-ing-function pointer on construction. Originally, the program had four 'act'-ing-functions defined: one that read from standard input from user as a human player, another that randomly placed pieces, Min-Max algorithm, and a Greedy algorithm. By changing the function pointers of

the objects, we could easily remove the human-player, and have two computer-players compete.

Now that we knew we had a solid base, we started to consider in more detail the implementation of our “learning” ‘act’ing-function. Even though theoretically any “learning” function would do, we felt we would have the best results with a typical neural net design. Then came the more difficult part where we had to decide what should be a part of the input layer and output layer, and how it should be represented.

At first we considered (the very silly idea of) somehow inputting the current board, and the available pieces, and then having the neural net output which piece, x-y coordinate, and rotation/flip for that piece. That idea quickly fell apart when attempting to implement a variable-length list of shapes, as well as teaching the neural net, without enforcing a ‘perfect piece’, that would be required for this implementation. The second idea was to input a future possible board, and have the neural net output a score on the quality of that choice, or how wise that choice is for the player. ‘Error’ in the system can be the difference between the neural network’s output compared to its final placement (rank). Because rank has limited resolution, and can therefore be difficult for a learning algorithm learn from, we choose to have the ‘rank’ equal to the difference of this player’s score, to the player with highest score, and the player with the highest score to be compared with the player with the second highest score. For example, if the final scores were {45, 56, 42, 69}, then the ‘rank’ would be respectively {-24, -13, -27, +13}.

calculate current game state, as well as . Such input the host-program can easily compute, and a output it can work with to make the best move, simply by picking the highest-scored future move. With this particular output of the neural net, training is very easy. To improve the neural net. we can compare the output of the neural net to the final score of the player at the end of the game. The representation of the future board can be given as a large set of integers, being that the blokus board is a set of small squares, which can only have 5 possible states. (1 for Player1, 2 for Player2, ..., 0 for empty space) This makes the ‘act’-ing function for our “learning” players equivalent to:

```
possible_moves = getFutureMoves(available_pieces)
possible_moves_scores = []
for move in possible_moves:
    future_board = current_board.copy()
    future_board.playMove(move);
    input_vector = [];
    for x in range(0, future_board.width):
        for y in range(0, future_board.height):
            cellvalue = 0;
            if(future_board.WhoOwns(x, y) == player1):
                cellvalue = 1;
            if(future_board.WhoOwns(x, y) == player2):
                cellvalue = 2;
            if(future_board.WhoOwns(x, y) == player3):
                cellvalue = 3;
            if(future_board.WhoOwns(x, y) == player4):
                cellvalue = 4;
```

```

        input_vector.append(cellvalue);
        score = evaluate_neural_network(weights, input_vector)
        possible_moves_scores.append(score)
    highest_score_index = # highest value's index in possible_moves_scores #
return possible_moves[highest_score_index]

```

Experimentation/Evaluation